

Multi-Layer Voice Reasoning Architecture for AI Systems (MLVRA)

Preserving Symbolic Fidelity Across the Voice AI Pipeline

Masoomul Haque Choudhury

Independent Researcher
Assam, India

June 2026

Abstract

Voice-based artificial intelligence systems exhibit a systematic and reproducible failure when confronted with tasks requiring exact symbolic reasoning -- tasks such as spelling verification, letter counting, rhyme identification, homophone disambiguation, and character-level pattern matching. This paper argues that the root cause is not a training data deficiency or a model capacity limitation, but a structural architectural flaw: the premature collapse of rich, multi-dimensional linguistic information into a single unified embedding space early in the inference pipeline. We propose the Multi-Layer Voice Reasoning Architecture (MLVRA), a framework that preserves four distinct but interconnected representational layers -- phonetic, canonical text, semantic, and confidence -- for every linguistic unit processed from speech input. A central component, the Dynamic Reasoning Router, identifies the reasoning modality demanded by each user query and selects the appropriate layer before reasoning begins. We describe the theoretical grounding of MLVRA in psycholinguistic research on parallel lexical access, detail the technical design of each layer and the router, analyze the failure modes of current pipeline architectures, and evaluate the potential impact of MLVRA across a range of high-stakes

application domains including medical dictation, legal transcription, and code generation. We further identify the principal open research challenges and propose a staged implementation roadmap.

Keywords: *voice AI, symbolic reasoning, phonetic representation, canonical text layer, dynamic routing, neuro-symbolic AI, ASR pipeline, multimodal cognition, spoken language understanding*

1. Introduction

The past decade has witnessed remarkable progress in voice-based artificial intelligence. Automatic speech recognition (ASR) systems have achieved near-human accuracy on standardized benchmark datasets, and large language models (LLMs) have demonstrated fluency and breadth of knowledge that was inconceivable a decade ago. Yet a peculiar and persistent failure mode remains underexplored in the research literature: voice AI systems consistently make errors on tasks that are trivially simple for human children.

Consider the following representative interaction. A user asks a voice assistant: "Which number below one thousand contains the letter A?" A correct system need only enumerate candidate numbers and check their spellings. Yet current voice AI systems frequently respond with answers like "forty" -- a number that contains no letter A -- and then, when prompted to reconsider, correct themselves. This is not a failure of general reasoning. The same system can solve multi-step arithmetic problems and synthesize information from diverse sources. The failure is specifically in the intersection of spoken language understanding and exact symbolic reasoning.

This paper investigates the structural cause of this failure and proposes a principled architectural solution. We argue that the failure arises because current voice AI pipelines collapse all information about a spoken word into a single representational format -- typically a dense semantic embedding -- at the earliest possible stage. This collapse is computationally efficient and works well for tasks that require semantic understanding. But it is fundamentally incompatible with tasks that require exact character-level reasoning, because the exact spelling structure is lost in the compression.

The Multi-Layer Voice Reasoning Architecture (MLVRA) proposed in this paper addresses this by maintaining parallel, separate representations of each linguistic unit throughout the processing pipeline. Rather than collapsing spoken input into a single embedding, MLVRA preserves a Phonetic Layer, a Canonical Text Layer, a Semantic Layer, and a Confidence and Ambiguity Layer simultaneously. A Dynamic Reasoning Router then

selects the appropriate layer based on the nature of each query, enabling exact symbolic reasoning when required without sacrificing performance on semantic tasks.

The contributions of this paper are as follows. First, we provide a detailed analysis of the architectural failure modes in current voice AI pipelines that produce symbolic reasoning errors. Second, we describe the psycholinguistic foundations of parallel representational processing in human language cognition. Third, we present the complete technical design of MLVRA, including all four layers and the Dynamic Reasoning Router. Fourth, we analyze the application domains where MLVRA would have the greatest practical impact. Fifth, we identify the principal open research challenges that must be addressed to realize MLVRA in production systems.

2. Background and Related Work

2.1 The Standard Voice AI Pipeline

Current voice AI systems follow a broadly consistent four-stage pipeline architecture. In Stage 1, raw audio waveforms are captured from microphone input. In Stage 2, an Automatic Speech Recognition (ASR) module converts the audio into text, typically producing a sequence of word tokens along with confidence scores. In Stage 3, the transcribed text is passed to a Large Language Model for understanding and response generation. In Stage 4, the generated text response is converted to audio output via a Text-to-Speech (TTS) module.

This pipeline has been optimized extensively for semantic fluency and conversational naturalness. The ASR stage has benefited from transformer-based architectures such as Whisper, achieving word error rates below 5% on standard benchmarks. The LLM stage has benefited from scale, instruction tuning, and reinforcement learning from human feedback (RLHF), producing outputs that are increasingly coherent and contextually appropriate.

However, the pipeline treats the ASR output as a purely semantic object -- a sequence of word tokens whose internal structure is not preserved downstream. Once the ASR module emits the token "forty," the downstream LLM has access only to the semantic identity of that token, not to its character-level structure. The information that "forty" is spelled f-o-r-t-y is not explicitly preserved in the token representation passed to the LLM.

2.2 Psycholinguistic Foundations

The MLVRA architecture is grounded in decades of psycholinguistic research on how humans process spoken language. The parallel distributed processing (PDP) models of McClelland and Rumelhart (1981) demonstrated that human word recognition involves simultaneous activation across multiple representational levels. The

Interactive Activation model proposed that phonological, orthographic, and semantic representations are activated in parallel and interact bi-directionally, rather than being processed strictly serially.

Coltheart's Dual Route Cascaded (DRC) model of reading provides additional theoretical grounding. The DRC model proposes that humans access word meaning through both a lexical route (whole-word recognition) and a non-lexical route (grapheme-to-phoneme conversion). The coexistence of these routes allows humans to handle both familiar words and novel letter strings, a flexibility that current voice AI systems lack.

Critically, experimental evidence from priming studies, N400 ERP components, and dual-task interference paradigms consistently demonstrates that orthographic (spelling) representations are activated automatically during spoken word recognition, even in tasks that do not require spelling awareness. This suggests that the separation of phonetic and orthographic information is not merely a convenience for reasoning -- it is a fundamental feature of human language cognition that enables flexible, task-appropriate access to word knowledge.

2.3 Related Architectural Proposals

Several lines of research in the AI community are relevant to MLVRA, though none directly addresses the specific failure mode of symbolic reasoning collapse in voice pipelines.

Neuro-symbolic AI research seeks to combine the pattern recognition capabilities of neural networks with the exact logical reasoning capabilities of symbolic systems. Systems such as Neural Theorem Provers and DeepMind's AlphaCode demonstrate that hybrid architectures can achieve performance that neither purely neural nor purely symbolic approaches can match. MLVRA can be understood as a neuro-symbolic approach specifically applied to the domain of spoken language understanding.

Mixture-of-Experts (MoE) architectures, such as those used in Mixtral and Switch Transformer, demonstrate that routing queries to specialized sub-networks based on their content is both feasible and effective at scale. The Dynamic Reasoning Router in MLVRA draws on this principle, though it operates at the level of representational layer selection rather than model weight selection.

Tool-augmented language models, including systems that can call Python interpreters or query databases, demonstrate that LLMs can be extended with exact symbolic computation capabilities when appropriate. MLVRA's Canonical Text Layer can be conceptualized as a form of built-in tool augmentation specifically for character-level operations.

3. Analysis of the Failure Mode

3.1 The Information Collapse Problem

The central failure mode addressed by MLVRA can be precisely characterized as premature information collapse. A spoken word carries at least three distinct types of information simultaneously: its acoustic-phonetic structure (how it sounds), its orthographic structure (how it is spelled), and its semantic content (what it means). These three types of information are related but not reducible to each other. The word "night" and the word "knight" are semantically distinct but phonetically identical. The word "read" has a single orthographic form but multiple phonetic realizations depending on tense. The word "queue" has a phonetically simple structure but an orthographically complex one.

Current voice AI pipelines collapse all three types of information into a single token representation at the ASR stage. The token "forty" that emerges from the ASR module is simply an index into a vocabulary -- it carries no explicit encoding of its phonetic structure, no explicit encoding of its orthographic structure (f-o-r-t-y), and its semantic information is available only implicitly through its embedding in the LLM's weight space.

When the downstream LLM is asked whether this token contains the letter "A," it must somehow recover the orthographic structure from the semantic embedding. This recovery is unreliable because the LLM's embedding space is optimized for semantic similarity, not orthographic fidelity. Semantically similar words ("forty," "fourteen," "fourth") may have similar embeddings despite having different orthographic structures, while orthographically similar words ("forty," "forte") may have dissimilar embeddings because their meanings differ.

3.2 Failure Taxonomy

We identify four distinct categories of task where premature information collapse produces systematic errors in current voice AI systems.

3.2.1 Letter-Level Analysis Tasks

Tasks requiring exact character-level analysis of word spellings are the most directly affected. Examples include letter counting ("How many letters are in the word you just said?"), letter presence checking ("Does that word contain the letter E?"), anagram detection ("Is that an anagram of another word?"), and spelling dictation ("Please spell that word aloud"). In each case, the task requires access to an explicit orthographic representation that current systems do not preserve.

3.2.2 Phonetic Reasoning Tasks

Tasks requiring exact phonetic analysis face a symmetric problem: the orthographic representation produced by ASR may obscure phonetic information. Rhyme detection is the canonical example. Two words rhyme if and only if their terminal phoneme sequences match, but this property is not reliably inferrable from spelling alone. "Though" and "go" rhyme; "though" and "through" do not. A system reasoning from orthography to phonetics will make systematic errors on such cases. Similarly, homophone identification requires access to phonetic representations that are not captured by the orthographic output of ASR.

3.2.3 Ambiguity Resolution Tasks

Speech is inherently ambiguous in ways that text is not. Current systems resolve such ambiguities immediately in the ASR stage, committing to a single interpretation and discarding the alternatives. This prevents downstream reasoning from taking ambiguity into account, leading to errors in contexts where the initial disambiguation was incorrect.

3.2.4 Cross-Modal Verification Tasks

A fourth category involves tasks where the user needs to verify that a spoken input matches a written form, or vice versa. Password and username dictation, code variable name verification, proper name spelling confirmation, and similar tasks require tight coordination between phonetic and orthographic representations. Without explicit preservation of both, such tasks cannot be reliably performed.

4. The Multi-Layer Voice Reasoning Architecture

MLVRA introduces a parallel representational structure that preserves distinct types of linguistic information throughout the voice AI pipeline, from speech input to final reasoning. Rather than collapsing information into a single format at the earliest opportunity, MLVRA maintains four layers simultaneously and routes reasoning queries to the appropriate layer via a Dynamic Reasoning Router.

4.1 Architectural Overview

At a high level, MLVRA replaces the single token representation produced by current ASR systems with a structured four-tuple for each processed linguistic unit. Each element of the tuple is a representation in a different format, maintained separately throughout the pipeline. The Dynamic Reasoning Router is responsible for selecting which element(s) of the tuple to use when answering a given user query. The architecture is designed to be modular and incrementally adoptable.

4.2 Layer 1: The Phonetic Layer

The Phonetic Layer preserves the acoustic-phonetic structure of each spoken word as a sequence of phonemes in International Phonetic Alphabet (IPA) notation, along with associated prosodic information including stress patterns, intonation contours, and speaking rate.

The Phonetic Layer serves as the primary resource for the following reasoning tasks: rhyme detection, where two words rhyme if and only if their phoneme sequences share a common terminal segment; homophone identification, where two words are homophones if their phoneme sequences are identical under standard pronunciation; accent and dialect normalization; and speech correction feedback.

Implementation of the Phonetic Layer can leverage existing phoneme recognition systems such as deep neural network acoustic models trained on phoneme-annotated speech corpora. The G2P (grapheme-to-phoneme) conversion systems used in TTS pipelines can also be inverted to generate phonetic representations from ASR hypotheses. The key design requirement is that the phonetic representation be stored explicitly rather than discarded after token generation.

4.3 Layer 2: The Canonical Text Layer

The Canonical Text Layer preserves the exact orthographic form of each spoken word as a sequence of individual characters. This layer is the primary resource for all character-level reasoning tasks and the most critical innovation of MLVRA from an engineering standpoint, because it represents the information most dramatically discarded by current pipeline architectures.

For the spoken word "forty," the Canonical Text Layer stores the explicit character sequence [f, o, r, t, y]. This representation directly enables letter presence checking (the character "a" is not in this sequence, so the answer is definitively "no"), letter counting (the sequence has five elements), anagram detection, alphabetical sorting, and pattern matching.

For spoken inputs that correspond to non-standard orthographic forms -- proper names, technical jargon, code identifiers, acronyms -- the Canonical Text Layer must handle uncertainty about the correct spelling. In such cases, the layer may store multiple candidate spellings with associated confidence scores, or flag the item for clarification.

4.4 Layer 3: The Semantic Layer

The Semantic Layer corresponds most closely to the representation currently produced by existing voice AI pipelines. It stores the meaning and contextual associations of each word as a dense embedding vector in a high-dimensional semantic space, of the kind produced by transformer-based language model encoders.

The Semantic Layer is the primary resource for the tasks at which current systems excel: conversational understanding, inference, question answering about world knowledge, summarization, and reasoning about meaning and implication. MLVRA does not diminish the importance of the Semantic Layer -- on the contrary, by offloading character-level and phonetic tasks to their appropriate layers, MLVRA allows the Semantic Layer to focus exclusively on tasks for which it is well-suited.

4.5 Layer 4: The Confidence and Ambiguity Layer

The Confidence and Ambiguity Layer is the most novel component of MLVRA from a theoretical standpoint. Current ASR systems produce not only their best-guess transcription but also alternative hypotheses with associated probability scores -- these are typically discarded before reaching the LLM. The Confidence and Ambiguity Layer preserves this uncertainty information explicitly and makes it available for downstream reasoning.

For a spoken input that the ASR module interprets as "cache" with 85% confidence and "cash" with 15% confidence, the Confidence and Ambiguity Layer stores the full distribution: {cache: 0.85, cash: 0.15}. This distribution is not collapsed to the most likely hypothesis but maintained as a structured object that can be reasoned about explicitly.

The Confidence and Ambiguity Layer enables several important capabilities: recognizing when uncertainty is too high to proceed confidently (triggering a clarification request), holding multiple hypotheses in parallel and revisiting reasoning if new context disambiguates the input, and providing uncertainty-aware reasoning where the system can report its confidence level alongside its answer.

4.6 The Dynamic Reasoning Router

The Dynamic Reasoning Router is the central coordinating mechanism of MLVRA. Its function is to analyze each incoming user query, identify the type of reasoning required to answer it, and direct that reasoning to the appropriate representational layer or combination of layers.

The router is implemented as a lightweight classification model that takes the user's query as input and produces a routing decision as output. In the simplest case, a query such as "Does it contain the letter A?" produces a single-layer routing decision to the Canonical Text Layer. A more complex query may produce a multi-layer routing decision engaging both the Phonetic Layer and the Semantic Layer in sequence.

Table 1: Query Type to Reasoning Layer Mapping

- Letter presence/counting -> Canonical Text Layer
- Spelling and pattern matching -> Canonical Text Layer
- Rhyme and homophone detection -> Phonetic Layer
- Meaning, inference, conversation -> Semantic Layer
- Disambiguation and clarification -> Confidence and Ambiguity Layer
- Code dictation and password verification -> Canonical Text + Confidence Layers
- Cross-modal verification -> Canonical Text + Phonetic Layers

The router's classification model can be trained on a relatively small labeled dataset of query-type pairs, since the number of routing categories is limited and the linguistic cues are often explicit. Queries containing words such as "spell," "letter," "character," or "contains" are strong signals for routing to the Canonical Text Layer. The router can be implemented with a few hundred parameters -- many orders of magnitude smaller than the LLM it coordinates -- adding negligible latency to the pipeline.

5. Theoretical Analysis

5.1 Information-Theoretic Perspective

From an information-theoretic standpoint, the failure of current systems can be understood as a form of lossy compression that discards information needed for a class of tasks. Shannon's source coding theorem establishes that any lossy compression scheme involves irreversible information loss. When the ASR pipeline maps a phoneme sequence to the token "forty," it performs a lossy transformation -- the phonetic detail is lost. When the LLM embedding space maps "forty" to a dense vector optimized for semantic similarity, it performs another lossy transformation -- the orthographic detail is effectively lost.

MLVRA addresses this by maintaining lossless representations at multiple levels of abstraction simultaneously. The cost of this approach is increased memory requirements -- four parallel representations rather than one -- but the benefit is that no information needed for any foreseeable reasoning task is discarded prematurely.

5.2 Computational Complexity

The computational overhead of MLVRA relative to current pipeline architectures is bounded and predictable. The Phonetic Layer requires running a phoneme recognition module on the audio input -- a task that modern acoustic models perform at roughly real-time speed. The Canonical Text Layer is already implicitly available as the direct output of the ASR module and requires no additional computation to store explicitly. The Semantic Layer corresponds to the existing LLM embedding computation. The Confidence and Ambiguity Layer requires storing the n-best list output that ASR modules already produce but currently discard.

The only genuinely novel computational requirement introduced by MLVRA is the Dynamic Reasoning Router, and as argued in Section 4.6, this can be implemented as an extremely lightweight model. The dominant computational costs of the pipeline -- ASR inference and LLM inference -- are unchanged. MLVRA is therefore a low-overhead architectural modification rather than a computationally expensive new approach.

5.3 Relationship to Human Cognitive Architecture

The parallel between MLVRA and human cognitive architecture is more than metaphorical. As reviewed in Section 2.2, psycholinguistic research has established that human word recognition involves genuinely parallel, independent activation of phonological, orthographic, and semantic representations. The MLVRA layers correspond to these independently established cognitive components.

This correspondence suggests that MLVRA is not merely a pragmatic engineering fix but an approximation of a computational architecture that evolution has converged upon for good functional reasons. Humans require flexible access to different types of linguistic information for different tasks, and maintaining these representations separately appears to be the solution that biological neural systems have developed. MLVRA proposes that artificial systems should adopt the same solution.

6. Application Domains and Impact Analysis

The failure modes addressed by MLVRA have different severity levels across application domains. In casual consumer voice assistant interactions, exact symbolic reasoning errors are rarely consequential. However, in a

range of high-stakes professional and specialized domains, such errors can have serious consequences. We analyze the most significant application domains below.

6.1 Medical Dictation

Medical dictation systems are perhaps the highest-priority application domain for MLVRA. The medical vocabulary is filled with phonetically similar but semantically distinct terms: ileum versus ilium, hypo- versus hyper-, perineal versus peroneal, temporal versus temporomandibular. A dictation error in a clinical note can affect patient care decisions with direct health consequences. Current voice-based medical dictation systems rely heavily on contextual disambiguation and domain-specific language models, but these cannot fully substitute for character-level precision when the phonetically similar terms appear in overlapping clinical contexts.

MLVRA's Canonical Text Layer and Confidence and Ambiguity Layer would directly address this failure mode. By preserving the phonetic ambiguity explicitly and flagging cases where two plausible interpretations have different character-level structures, MLVRA would enable the system to request clarification precisely when needed -- reducing both the error rate and the rate of unnecessary interruptions.

6.2 Legal Transcription

Legal transcription requires word-for-word accuracy because the transcribed text may become a matter of legal record. Voice-based transcription systems used in courtrooms, depositions, and contract dictation need to distinguish between "shall" and "should," "may" and "must," "the defendant" and "a defendant." These distinctions are often phonetically subtle and semantically significant. Beyond individual word accuracy, legal transcription benefits from MLVRA's ability to handle proper names accurately.

6.3 Code Dictation and Programming Assistance

Voice-based coding assistance is an emerging use case that places particularly stringent requirements on character-level accuracy. Variable names, function names, and string literals must be transcribed exactly, since a single character difference produces a syntax error or semantic bug. A spoken variable name like "userAuthToken" must be preserved as that exact sequence of characters, not normalized to a semantically equivalent synonym. MLVRA's Canonical Text Layer would directly support code dictation by maintaining the exact character sequence of each spoken identifier.

6.4 Accessibility Technology

Voice-based interfaces are particularly important for users with motor disabilities, visual impairments, or other conditions that make traditional keyboard and touchscreen interfaces difficult or impossible. For such users, the accuracy and reliability of voice AI systems are not merely conveniences but requirements for independent functioning. MLVRA's improvements to symbolic reasoning accuracy would have a disproportionate positive impact on this population.

6.5 Multilingual and Cross-Lingual Applications

Voice AI systems used in multilingual contexts face compounded challenges: the same phoneme sequence may correspond to words in different languages with different spellings and meanings. MLVRA's modular layer architecture is particularly well-suited to multilingual applications. Each layer can be instantiated with language-specific modules -- a multilingual phoneme recognizer for the Phonetic Layer, a multilingual character tokenizer for the Canonical Text Layer -- while the Dynamic Reasoning Router operates at a language-agnostic level of abstraction.

6.6 Military and Emergency Communication

High-stakes communication in military and emergency response contexts requires precision under adverse acoustic conditions. Callsigns, map coordinates, equipment identifiers, and medical status codes must be transmitted and received with exact fidelity. The NATO phonetic alphabet (Alpha, Bravo, Charlie...) exists precisely because letter-level precision is critical in such contexts. Voice AI systems deployed in these contexts would benefit substantially from MLVRA's preservation of phonetic and orthographic representations.

7. Implementation Roadmap and Technical Challenges

7.1 Staged Implementation

We propose a three-stage implementation roadmap for MLVRA. Each stage delivers tangible benefits and can be deployed independently, allowing incremental adoption.

Stage 1: Layer Preservation (6-12 months)

The first and most impactful stage is the simplest: modify existing ASR pipelines to preserve all four types of representation rather than discarding them after token generation. This requires no changes to the LLM component and no new training. The ASR module already computes phoneme probabilities and n-best lists internally -- Stage 1 simply ensures these are stored and passed downstream rather than discarded. This stage alone would address a significant fraction of the failure cases described in Section 3.

Stage 2: Dynamic Routing (12-24 months)

Stage 2 implements the Dynamic Reasoning Router as a lightweight classification model trained on a dataset of query-routing pairs. This dataset can be generated synthetically from templates or collected through active learning, starting with the clearest routing cases and progressively adding ambiguous cases as the router's confidence calibration improves. Initial implementations may target response times under 10 milliseconds, well within the acceptable latency budget for voice interaction.

Stage 3: Integrated Reasoning (24-48 months)

Stage 3 involves deeper integration of the multi-layer representations into the LLM's reasoning process. Rather than providing the LLM with a single text input, Stage 3 systems provide structured inputs that explicitly encode all four layers. This may require fine-tuning the LLM on multi-layer input formats, or designing new attention mechanisms that can differentially weight different representational layers.

7.2 Open Research Challenges

Several significant research challenges must be addressed to realize MLVRA at full scale:

- **Router Ambiguity Resolution:** Many real-world queries require multiple layers simultaneously or involve reasoning that spans layers. Developing principled methods for layer combination and query decomposition is an open research problem.
- **Unknown Vocabulary Handling:** The Canonical Text Layer faces its greatest challenge with spoken proper names, technical terms, and novel coinages that do not appear in standard vocabulary databases. Zero-shot orthographic inference from phonetics alone remains an unsolved problem.
- **Latency Optimization:** The Phonetic Layer requires running a full phoneme recognition model on the audio input, which may add perceptible latency in streaming (real-time) applications. Efficient streaming architectures for phoneme recognition are needed.
- **Confidence Calibration:** The Confidence and Ambiguity Layer's usefulness depends on the ASR module producing well-calibrated confidence scores. Current ASR systems are often overconfident; improving confidence calibration is an active area of research.
- **Cross-Layer Consistency:** It is possible for the four layers to produce inconsistent information about a single spoken word. Resolving such inconsistencies in a principled way requires new algorithms for multi-layer agreement and conflict resolution.

8. Discussion

8.1 Limitations of the Proposed Architecture

MLVRA addresses the premature information collapse problem but does not solve all the challenges of voice AI. The architecture assumes that the four proposed layers are sufficient to capture all the information needed for all voice AI tasks. This assumption may be incorrect. Future work may identify additional types of information -- pragmatic intent, discourse structure, paralinguistic cues such as emotion and emphasis -- that current systems also discard prematurely.

The Dynamic Reasoning Router introduces a classification step that is itself a potential failure point. If the router incorrectly classifies a query's reasoning requirements, the resulting error may be harder to diagnose than the original failure mode it was designed to prevent. Router failure modes must be carefully characterized and mitigated.

Finally, MLVRA in its current formulation addresses the single-utterance case -- a single spoken word or phrase whose representations are maintained through the pipeline. Extending the architecture to handle multi-word utterances where different words require different layer combinations at different points in the reasoning process is a significant additional complexity that this paper does not fully address.

8.2 Broader Implications

The analysis presented in this paper suggests a broader principle that extends beyond voice AI: AI systems should be designed to preserve the full dimensionality of information about their inputs rather than compressing that information to the minimum required for the most common tasks. The efficiency gains from early compression are real but may be outweighed by the loss of capability for tasks that depend on the compressed information.

This principle has implications for AI system design beyond the voice modality. Text-based systems that compress input text to semantic embeddings before reasoning may similarly lose information needed for certain tasks. Image processing systems that discard fine-grained pixel-level information in favor of feature-level representations may perform poorly on tasks that require exact pixel-level analysis. The general lesson of MLVRA is that premature information compression is a design choice with hidden costs that should be made explicit and reconsidered.

9. Conclusion

This paper has presented the Multi-Layer Voice Reasoning Architecture (MLVRA), a principled architectural solution to the systematic failure of current voice AI systems on tasks requiring exact symbolic reasoning. The

core insight is that current systems fail not because they lack the capability for symbolic reasoning, but because they discard the information needed for such reasoning at the earliest stage of the pipeline.

MLVRA proposes maintaining four parallel representational layers -- phonetic, canonical text, semantic, and confidence -- for each linguistic unit processed from speech input, and routing reasoning queries to the appropriate layer via a Dynamic Reasoning Router. This architecture is grounded in psycholinguistic research on parallel representational processing in human cognition, consistent with emerging directions in neuro-symbolic AI and mixture-of-experts research, and amenable to staged implementation that delivers incremental benefits.

The impact of MLVRA would be most significant in high-stakes application domains -- medical dictation, legal transcription, code dictation, accessibility technology -- where the symbolic reasoning errors of current systems have real consequences. But the broader implication of this work is a principle of information preservation that should inform AI system design across modalities and application domains: do not compress what you may need to reason about.

The open research challenges identified in this paper define a research agenda for the voice AI community. We hope this paper contributes to motivating and orienting that agenda, and that the MLVRA framework provides a useful conceptual foundation for future work in symbolic-aware spoken language understanding.

References

- [1] Coltheart, M., Rastle, K., Perry, C., Langdon, R., and Ziegler, J. (2001). DRC: A Dual Route Cascaded model of visual word recognition and reading aloud. *Psychological Review*, 108(1), 204-256.
- [2] McClelland, J. L., and Rumelhart, D. E. (1981). An interactive activation model of context effects in letter perception. *Psychological Review*, 88(5), 375-407.
- [3] Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., and Sutskever, I. (2023). Robust Speech Recognition via Large-Scale Weak Supervision. *Proceedings of the 40th International Conference on Machine Learning (ICML 2023)*.

- [4] Shazeer, N., Mirhoseini, A., Maziarz, K., Davis, A., Le, Q., Hinton, G., and Dean, J. (2017). Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of-Experts Layer. ICLR 2017.
- [5] Lake, B. M., Ullman, T. D., Tenenbaum, J. B., and Gershman, S. J. (2017). Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40, e253.
- [6] Marcus, G. (2019). *Rebooting AI: Building Artificial Intelligence We Can Trust*. Pantheon Books.
- [7] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention Is All You Need. NeurIPS 2017.
- [8] Jiang, A. Q., et al. (2024). Mixtral of Experts. arXiv:2401.04088.
- [9] Kuhl, P. K. (2004). Early language acquisition: Cracking the speech code. *Nature Reviews Neuroscience*, 5(11), 831-843.
- [10] Harm, M. W., and Seidenberg, M. S. (2004). Computing the meanings of words in reading: Cooperative division of labor between visual and phonological processes. *Psychological Review*, 111(3), 662-720.
- [11] Brown, T. B., et al. (2020). Language Models are Few-Shot Learners. NeurIPS 2020.
- [12] Bojanowski, P., Grave, E., Joulin, A., and Mikolov, T. (2017). Enriching Word Vectors with Subword Information. *Transactions of the ACL*, 5, 135-146.
- [13] Hochreiter, S., and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8), 1735-1780.
- [14] Shannon, C. E. (1948). A Mathematical Theory of Communication. *Bell System Technical Journal*, 27(3), 379-423.
- [15] Minsky, M. (1986). *The Society of Mind*. Simon and Schuster.
-

Acknowledgements

The author is grateful to the broader research communities in speech processing, natural language understanding, and cognitive science whose foundational work provides the scaffolding for the ideas presented here. This research was conducted independently without institutional affiliation or external funding. The author thanks the open-source research community for making foundational papers and technical resources accessible to independent researchers worldwide.

Masoomul Haque Choudhury | Independent Researcher | Assam, India | June 2026